

ata-analysis-and-income-prediction

November 12, 2023

```
[1]: # This Python 3 environment comes with many helpful analytics libraries
      ↳ installed
      # It is defined by the kaggle/python Docker image: https://github.com/kaggle/
      ↳ docker-python
      # For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import math

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list
↳ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that
↳ gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved
↳ outside of the current session
```

/kaggle/input/top-100-popular-movies-from-2003-to-2022-imdb/movies.csv

1 Business Data Challenge

Companies in the entertainment industry nowadays face a substantial risk: Large upfront investments are often necessary for movie productions, but only 20% of movies end up being profitable (Forbes, 2019). Indicators are needed that help companies estimate the financial of their movie, make more informed decisions and so are able to mitigate the risk of losing money on a movie project. To enhance the movie production company's chances of producing profitable films the use of data analytics is crucial. By analyzing historical data on movie success, trends, correlations, and patterns can be uncovered and used to make informed decisions regarding genre, casting, or marketing strategies. Utilizing machine learning algorithms, our predictive models will be able to forecast potential box office performance and allow to optimize resources and achieve maximum

profitability.

2 Load Data Set

```
[2]: df = pd.read_csv('/kaggle/input/top-100-popular-movies-from-2003-to-2022-imdb/
↳movies.csv')
```

3 Inspect Data

```
[3]: # overview
df.head()
```

```
[3]:
```

	Title	Rating	Year	Month	Certificate	Runtime \
0	Avatar: The Way of Water	7.8	2022	December	PG-13	192
1	Guillermo del Toro's Pinocchio	7.6	2022	December	PG	117
2	Bullet Train	7.3	2022	August	R	127
3	The Banshees of Inisherin	7.8	2022	November	R	114
4	M3gan	6.4	2022	December	PG-13	102

	Directors \
0	James Cameron
1	Guillermo del Toro, Mark Gustafson
2	David Leitch
3	Martin McDonagh
4	Gerard Johnstone

	Stars \
0	Sam Worthington, Zoe Saldana, Sigourney Weaver...
1	Ewan McGregor, David Bradley, Gregory Mann, Bu...
2	Brad Pitt, Joey King, Aaron Taylor Johnson, Br...
3	Colin Farrell, Brendan Gleeson, Kerry Condon, ...
4	Jenna Davis, Amie Donald, Allison Williams, Vi...

	Genre	Filming_location	Budget \
0	Action, Adventure, Fantasy	New Zealand	\$350,000,000
1	Animation, Drama, Family	USA	\$35,000,000
2	Action, Comedy, Thriller	Japan	\$85,900,000
3	Comedy, Drama	Ireland	Unknown
4	Horror, Sci-Fi, Thriller	New Zealand	\$12,000,000

	Income	Country_of_origin
0	\$2,267,946,983	United States
1	\$108,967	United States, Mexico, France
2	\$239,268,602	Japan, United States
3	\$19,720,823	Ireland, United Kingdom, United States
4	\$171,253,910	United States

```
[4]: # overview
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Title                 2000 non-null   object
 1   Rating               1999 non-null   float64
 2   Year                 2000 non-null   int64
 3   Month               2000 non-null   object
 4   Certificate          1966 non-null   object
 5   Runtime              2000 non-null   object
 6   Directors            2000 non-null   object
 7   Stars                2000 non-null   object
 8   Genre                2000 non-null   object
 9   Filming_location     2000 non-null   object
10   Budget               2000 non-null   object
11   Income               2000 non-null   object
12   Country_of_origin    2000 non-null   object
dtypes: float64(1), int64(1), object(11)
memory usage: 203.2+ KB
```

```
[5]: # list all columns
df.columns
```

```
[5]: Index(['Title', 'Rating', 'Year', 'Month', 'Certificate', 'Runtime',
          'Directors', 'Stars', 'Genre', 'Filming_location', 'Budget', 'Income',
          'Country_of_origin'],
          dtype='object')
```

```
[6]: # check duplicates
df.loc[df.duplicated(keep=False)]
```

```
[6]: Empty DataFrame
Columns: [Title, Rating, Year, Month, Certificate, Runtime, Directors, Stars,
Genre, Filming_location, Budget, Income, Country_of_origin]
Index: []
```

```
[7]: # check unique values
df.nunique()
```

```
[7]: Title                1989
Rating                 62
Year                   20
Month                  14
```

```

Certificate      12
Runtime          113
Directors        1082
Stars            1990
Genre            244
Filming_location 97
Budget           305
Income           1856
Country_of_origin 406
dtype: int64

```

```
[8]: df['Budget'].unique()
```

```

[8]: array(['$350,000,000 ', '$35,000,000 ', '$85,900,000 ', 'Unknown',
'$12,000,000 ', '$120,000,000 ', '$80,000,000 ', '$20,000,000 ',
'$10,000,000 ', '$40,000,000 ', '$78,000,000 ', '$1,000,000 ',
'$16,000,000 ', '$195,000,000 ', '$100,000,000 ', '$250,000,000 ',
'€ 10,000,000', '$25,000,000 ', '$3,000,000 ', '$150,000,000 ',
'$170,000,000 ', '$17,000,000 ', '$200,000,000 ', '$4,500,000 ',
'$32,000,000 ', '$60,000,000 ', '$72,000,000 ', '$2,000,000 ',
'$68,000,000 ', '$24,000,000 ', '$9,000,000 ', '$50,000,000 ',
'$90,000,000 ', '$250,000 ', '$85,000,000 ', '$8,000,000 ',
' 3,500,000,000', '$22,000,000 ', '€ 7,500,000', '$165,000,000 ',
'CA$15,000', '$30,000,000 ', '$185,000,000 ', '$75,000,000 ',
'$18,000,000 ', '$55,000,000 ', '$43,000,000 ', '€ 5,000,000',
'$300,000,000 ', 'SEK\xa019,000,000', '$190,000,000 ',
'$15,000,000 ', '$160,000,000 ', '$70,000,000 ', '$24,350,000 ',
'$820,000 ', '$110,000,000 ', '$1,100,000 ', '€ 5,700,000',
'$39,000,000 ', '$205,000,000 ', '$6,000,000 ', '$65,000,000 ',
'€ 14,000,000', '$14,000,000 ', '$84,500,000 ', '$45,000,000 ',
'$67,000,000 ', '$5,000,000 ', '$50,300,000 ', '$7,000,000 ',
'$34,000,000 ', '$135,000,000 ', '$33,000,000 ', '$3,500,000 ',
'$5,300,000 ', '$175,000,000 ', '$356,000,000 ', '$11,400,000 ',
'$95,000,000 ', '$159,000,000 ', '$275,000,000 ', '$97,600,000 ',
'$19,000,000 ', '€ 14,750,000', '$11,000,000 ', '$183,000,000 ',
'$260,000,000 ', '$125,000,000 ', '$18,600,000 ', '$79,000,000 ',
'€ 30,690,000', '$13,000,000 ', 'DKK\xa019,500,000',
'$26,000,000 ', '$48,000,000 ', '$115,000,000 ', '$5,400,000 ',
'$42,000,000 ', '$4,000,000 ', '$6,200,000 ', '$321,000,000 ',
'$23,000,000 ', '$69,000,000 ', '$36,000,000 ', '$178,000,000 ',
'$52,000,000 ', '€ 2,600,000', '$130,000,000 ', '$5,500,000 ',
'$1,700,000 ', '€ 8,700,000', '$37,000,000 ', '$29,000,000 ',
'$62,000,000 ', '$94,000,000 ', '$3,200,000 ', '$88,000,000 ',
'$8,500,000 ', '$59,000,000 ', '€ 8,000,000', '$162,000,000 ',
'$11,000 ', '$38,000,000 ', '$880,000 ', '€ 11', '$7,700,000 ',
'$28,000,000 ', '€ 4,000,000', '$97,000,000 ', '$217,000,000 ',
'$19,400,000 ', '$149,000,000 ', '$84,000,000 ', '$317,000,000 ',

```

```
'$180,000,000 ', '$230,000,000 ', '$58,000,000 ', '$104,000,000 ',
'$177,200,000 ', '£100,000', '£5,400,000', '$4,800,000 ',
'₩10,000,000,000', '$35,000 ', '$22,500,000 ', '$47,000,000 ',
'¥370,000,000', '$44,000,000 ', '$144,000,000 ', '$9,900,000 ',
'€ 3,500,000', '$46,000,000 ', '$140,000,000 ', '£500,000',
'$6,420,000 ', '€ 12,000,000', '$145,000,000 ', '$245,000,000 ',
'$108,000,000 ', '$31,000,000 ', '$11,800,000 ', '$14,800,000 ',
'$61,000,000 ', '$1,800,000 ', '$2,500,000 ', '$155,000,000 ',
'$9,400,000 ', '$74,000,000 ', '$50,100,000 ', '$53,000,000 ',
'$500,000 ', '$105,000,000 ', '$18,026,148 ', '$176,000,000 ',
'€ 6,000,000', '$8,900,000 ', '€ 12,300,000', '$49,000,000 ',
'$100,000 ', '$16,200,000 ', '$3,300,000 ', '$81,000,000 ',
'$210,000,000 ', '$127,000,000 ', '$58,800,000 ', '$5,100,000 ',
'$2,400,000 ', '$1,600,000 ', '$12,600,000 ', '$6,500,000 ',
'$19,800,000 ', '$66,000,000 ', '$13,300,000 ', '$225,000,000 ',
'$4,700,000 ', '$50,000 ', '$39,200,000 ', '$92,000,000 ',
'€ 4,645,437', '$215,000,000 ', '$76,000,000 ', '€ 9,200,000',
'$27,220,000 ', '€ 13,500,000', '€ 300', '$3,800,000 ',
'$220,000,000 ', '$21,000,000 ', '£3,000,000', '$270,000 ',
'$102,000,000 ', '$44,500,000 ', '$40,600,000 ', '$209,000,000 ',
'$12,500,000 ', '$750,000 ', '$242,000 ', 'A$3,000,000',
'€ 9,500,000', '$32,500,000 ', '$27,000,000 ', '$82,000,000 ',
'€ 10,002,914', '€ 1,500,000', '£20,000,000', '$7,400,000 ',
'$56,000,000 ', '$93,000,000 ', '$163,000,000 ', '$50,200,000 ',
'$300,000 ', '£3,500,000', '$2,600,000 ', 'NOK\xa019,900,000',
'£150,000,000', '$117,000,000 ', '€ 1,948,000', '$6,800,000 ',
'$1,500,000 ', '$10 ', '€ 4,830,000', '$10,500,000 ',
'$237,000,000 ', '$7,500,000 ', ' 550,000,000', '$23,600,000 ',
'$3,700,000 ', '$37,500,000 ', '£1,000,000', 'A$8,240,000',
'€ 3,390,000', '$15,500,000 ', '$9,750,000 ', '$230,000 ',
'$11,715,578 ', '$17,500,000 ', '£13,500,000', '£8,000,000',
'$258,000,000 ', '£25,000,000', '£26,000,000', '$200,000 ',
'$15,000 ', '$10,200,000 ', '€ 3,400,000', '$150,000 ',
'€ 2,400,000', '$270,000,000 ', '$72,500,000 ', '$82,500,000 ',
'£1,500,000', '$16,500,000 ', '£9,800,000', '₩12,215,500,000',
'¥300,000,000', '$54,000,000 ', '$207,000,000 ', 'CA$2,200,000',
'$113,000,000 ', '$132,000,000 ', '€ 18,151,814', '$126,000,000 ',
'$475,000 ', '$950,000 ', '$7,900,000 ', '₩4,200,000,000',
'$400,000 ', '$1,200,000 ', '£4,000,000', '£695,393',
'$2,700,000 ', 'CN¥100,000,000', '$7,000 ', '$8,600,000 ',
'£1,700,000', '$56,600,000 ', '$2,800,000 ', '$137,000,000 ',
'$109,000,000 ', '€ 2,200,000', '$9,500,000 ', '$900,000 ',
'$128,000,000 ', '€ 4,800,000'], dtype=object)
```

```
[9]: df['Runtime'].unique()
```

```
[9]: array(['192', '117', '127', '114', '102', '132', '134', '112', '151',
        '107', '188', '105', '131', '125', '126', '98', '101', '161',
        '140', '147', '139', '97', '154', '123', '130', '115', '176', '87',
        '129', '104', '99', '93', '137', '128', '100', '148', '158', '96',
        '86', '108', '124', '89', '111', '103', '135', '95', '119', '136',
        '77', '118', '138', '159', '122', '187', '146', '121', '91', '116',
        '167', '113', '110', '141', '92', '155', '163', '106', '133', '88',
        '150', '156', '242', '109', '143', '152', '83', '120', '90', '144',
        '179', '142', '160', '94', '85', '181', '209', '169', '84', '149',
        '78', '75', '164', '70', '145', '79', '82', '168', '165',
        'Unknown', '153', '180', '50', '157', '172', '80', '162', '170',
        '166', '81', '191', '71', '175', '201', '178'], dtype=object)
```

```
[10]: df['Certificate'].unique()
```

```
[10]: array(['PG-13', 'PG', 'R', 'TV-14', 'TV-MA', 'TV-PG', 'TV-Y7',
        'Not Rated', nan, 'NC-17', 'TV-G', 'Unrated', 'G'], dtype=object)
```

4 Data Cleaning

```
[11]: df.drop(['Title'], axis=1, inplace=True)
```

```
[12]: # percentage of nan values by column
df.isna().sum()
```

```
[12]: Rating          1
      Year            0
      Month           0
      Certificate      34
      Runtime         0
      Directors        0
      Stars            0
      Genre            0
      Filming_location 0
      Budget           0
      Income           0
      Country_of_origin 0
      dtype: int64
```

```
[13]: # updates the 'Certificate' value for the selected rows
# with the corresponding mode value associated with the first genre
invalid_certificates = ['Unrated', 'Not Rated', np.nan]
genre_mode = df.groupby(df['Genre'].str.split(',').str[0])['Certificate'].
    .transform(lambda x: x.mode().iloc[0])
mask = df['Certificate'].isin(invalid_certificates)
df.loc[mask, 'Certificate'] = genre_mode[mask]
```

```
[14]: # percentage of missing values recorded as "unknown"
def check_unknowns():
    for col in df.columns:
        unknown_percent = (df[col] == 'Unknown').sum() / len(df) * 100
        print(f'{col}: {unknown_percent:.2f}%')
    check_unknowns()
```

```
Rating: 0.00%
Year: 0.00%
Month: 0.00%
Certificate: 0.00%
Runtime: 0.10%
Directors: 0.00%
Stars: 0.00%
Genre: 0.00%
Filming_location: 3.80%
Budget: 15.20%
Income: 7.25%
Country_of_origin: 0.00%
```

```
[15]: # turn all unknown values into 0s for later processing
for col in ['Budget', 'Income', 'Runtime']:
    df[col].replace('Unknown', '0', inplace=True)
```

```
[16]: # conversion values 03/11/23
conversion_dict = {
    '$': 1.0,
    '€': 1.07,
    '£': 1.21,
    '₩': 0.00078,
    'C': 0.73,
    ' ': 0.012,
    '¥': 0.0075,
    'A': 0.68,
    'S': 0.095,
    'D': 0.14,
    'N': 0.10
}
```

```
[17]: import re
# strip each value of its symbol and convert the value
# based on the symbol as per the dictionary
for col in ['Budget', 'Income']:
    df[col] = df[col].apply(lambda x:
                            int(re.sub('[^0-9]', '', x)) *
                            ↪conversion_dict.get(x[0], 1))
```

```
[18]: df['Runtime'] = df['Runtime'].astype(int)
# turn all unknown values, 0s, into nans
df.loc[:, ['Budget',
           'Income',
           'Runtime']] = df.loc[:, ['Budget',
                                   'Income',
                                   'Runtime']].replace(0, np.nan)
```

```
[19]: # fill nan values for runtime, budget, income,
df['Runtime'].dropna(inplace=True)
df['Income'].fillna(df['Income'].median(),inplace=True)
df['Budget'].fillna(df['Budget'].median(),inplace=True)
```

```
[20]: df.isna().sum()
```

```
[20]: Rating          1
Year                0
Month              0
Certificate         0
Runtime            2
Directors          0
Stars              0
Genre              0
Filming_location   0
Budget             0
Income             0
Country_of_origin  0
dtype: int64
```

```
[21]: df.loc[df['Runtime'].isnull(), :]
```

```
[21]:
```

	Rating	Year	Month	Certificate	Runtime	Directors	\
848	2.1	2014	2014	R	NaN	Enzo Zelocchi	
1249	3.2	2010	September	PG-13	NaN	Enzo Zelocchi	

	Stars	\
848	Enzo Zelocchi, Miryam Negrin, Paul Gregory, Ma...	
1249	Enzo Zelocchi, Charlotte Labadie, David M Edel...	

	Genre	Filming_location	Budget	Income	\
848	Comedy, Drama, Thriller	USA	250000.0	90842646.0	
1249	Romance, Drama, Family	Unknown	10.0	90842646.0	

	Country_of_origin
848	United States
1249	United States


```
[22]: df.loc[df['Rating'].isnull(), :]
```

```
[22]:      Rating  Year  Month Certificate  Runtime  Directors \
85      NaN  2022  January      PG-13    126.0  Marc Forster

      Stars      Genre \
85  Tom Hanks, Rachel Keller, Manuel Garcia Rulfo,...  Comedy, Drama

      Filming_location  Budget  Income  Country_of_origin
85              USA  40000000.0  90842646.0  Sweden, United States
```

```
[23]: df.drop(index=[85, 1249, 848], inplace=True)

df.reset_index(drop=True, inplace=True)
```

```
[24]: # fill unknown values with most common
df['Filming_location'] = df['Filming_location'].replace('Unknown', 'USA')
```

```
[25]: # drop the two rows with wrong month value
df = df.drop(df[df['Month'] == '2014'].index)
df = df.drop(df[df['Month'] == '2008'].index)
```

```
[26]: # years from string to integers
df['Year'] = df['Year'].astype(int)
```

```
[27]: # check if all unknowns have been taken care of
check_unknowns()
```

```
Rating: 0.00%
Year: 0.00%
Month: 0.00%
Certificate: 0.00%
Runtime: 0.00%
Directors: 0.00%
Stars: 0.00%
Genre: 0.00%
Filming_location: 0.00%
Budget: 0.00%
Income: 0.00%
Country_of_origin: 0.00%
```

```
[28]: # check dataframe after cleaning
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1996 entries, 0 to 1996
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
#   ...
```

```

---  -----
0   Rating                1996 non-null    float64
1   Year                  1996 non-null    int64
2   Month                 1996 non-null    object
3   Certificate            1996 non-null    object
4   Runtime               1996 non-null    float64
5   Directors             1996 non-null    object
6   Stars                 1996 non-null    object
7   Genre                 1996 non-null    object
8   Filming_location      1996 non-null    object
9   Budget                1996 non-null    float64
10  Income                1996 non-null    float64
11  Country_of_origin     1996 non-null    object
dtypes: float64(4), int64(1), object(7)
memory usage: 202.7+ KB

```

5 Exploratory Data Analysis

```
[29]: df.describe().T
```

```

[29]:
count      mean      std      min      25%  \
Rating    1996.0  6.665932e+00  9.030395e-01    1.90  6.175000e+00
Year      1996.0  2.012498e+03  5.768334e+00  2003.00  2.007000e+03
Runtime   1996.0  1.132179e+02  1.961313e+01    50.00  9.900000e+01
Budget    1996.0  5.804498e+07  5.647101e+07    11.77  2.000000e+07
Income    1996.0  1.818519e+08  2.685744e+08   305.00  2.956873e+07

      50%      75%      max
Rating    6.7  7.300000e+00  9.000000e+00
Year     2012.5  2.017250e+03  2.022000e+03
Runtime   110.0  1.240000e+02  2.420000e+02
Budget   40000000.0  7.500000e+07  3.560000e+08
Income  90842646.0  2.084946e+08  2.922918e+09

```

```

[30]: # create a Profit column
df['Profit'] = (df['Income'] - df['Budget']).apply(lambda x: round(x, -6))

# Create a new column for ROI
df['ROI'] = (df['Income'] - df['Budget']) / df['Budget']

```

```
[31]: df['ROI'].describe()
```

```

[31]: count      1996.000000
mean         242.325442
std          7316.282702
min           -0.999992
25%           0.164218

```

```
50%          1.384460
75%          3.353567
max         282997.897196
Name: ROI, dtype: float64
```

```
[32]: # remove Budget values less than 1000 as those are likely typos or extreme
      ↪ outliers at best
      df = df[df['Budget'] > 1000]
```

```
[33]: import matplotlib.pyplot as plt
      import seaborn as sns

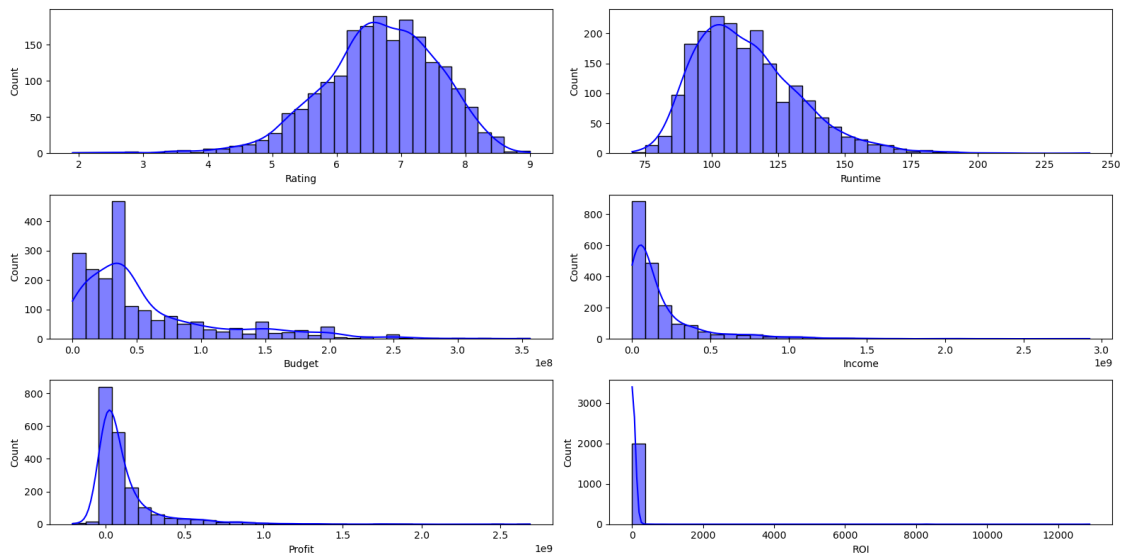
      # plot histogram for all numerical values
      num_values = ['Rating', 'Runtime', 'Budget', 'Income', 'Profit', 'ROI']

      fig, axs = plt.subplots(nrows=3, ncols=2, figsize=(16, 8))

      for i, col in enumerate(num_values):
          sns.histplot(data=df,
                       x=col,
                       bins=35,
                       kde=True,
                       color='blue',
                       ax=axs[i//2, i%2])
          axs[i//2, i%2].set_xlabel(col)

      # axs[2,1].axis('off')


      plt.tight_layout()
      plt.show()
```



```
[34]: fig, axs = plt.subplots(nrows=1, ncols=len(num_values), figsize=(16, 4))


for i, col in enumerate(num_values):
    sns.boxplot(x=df[col], ax=axs[i])
    axs[i].set_xlabel(col)

plt.tight_layout()
plt.show()
```



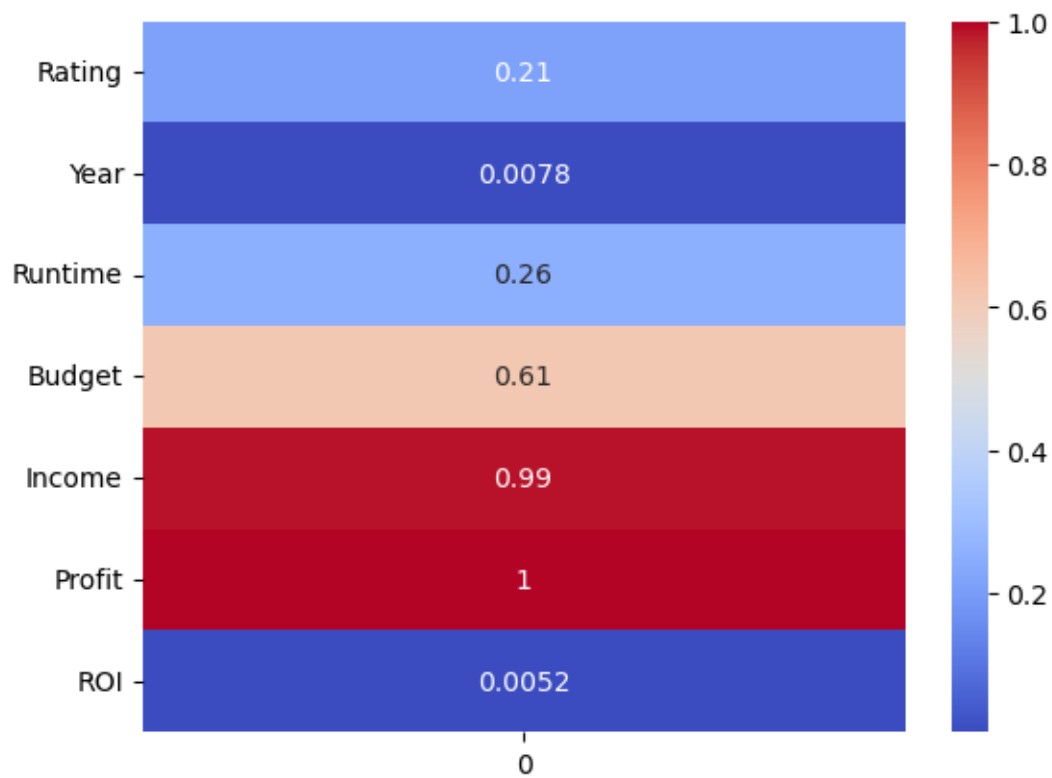
```
[35]: # Draw a categorical scatterplot to show each observation
ax = sns.scatterplot(data=df, x='Runtime', y='Profit')
ax.set(ylabel='Profit $bn')
ax.set(xlabel='Runtime (min)')
ax.set_xlim(60, 220)
plt.show
```

```
[35]: <function matplotlib.pyplot.show(close=None, block=None)>
```




```
[36]: # calculate the correlation coefficients
corr_matrix = df.corrwith(df['Profit'])

# plot the correlation matrix as a heatmap
sns.heatmap(corr_matrix.to_frame(), cmap='coolwarm', annot=True)
plt.show()
```



```
[37]: # create a scatter plot with a regression line
sns.lmplot(x='Budget', y='Profit', data=df, line_kws={'color': 'red'})
plt.title('Pearson Correlation')
plt.show()
```



```
[38]: def most_freq(col, n):
      tmp_df = df[[col]].copy()
      tmp_df[col] = [x.split(',') for x in tmp_df[col]]
      most_freq_elements = tmp_df[col].explode().value_counts().index[:n].tolist()
      return most_freq_elements
```

```
[39]: top_genre = most_freq('Genre', 10)
      top_stars = most_freq('Stars', 15)
      top_directors = most_freq('Directors', 10)
      top_country_of_origin = most_freq('Country_of_origin', 10)
      top_filming_location = most_freq('Filming_location', 10)
```

```
[40]: top_elements_dict = {'Stars': top_stars, 'Directors': top_directors,
                          'Genre': top_genre, 'Country_of_origin': top_country_of_origin,
                          'Filming_location': top_filming_location}
```

```
[41]: print('### Top 10 Genre ###\n ', top_genre)
      print('\n### Top 15 Stars ###\n ', top_stars)
      print('\n### Top 10 Directors ###\n ', top_directors)
      print('\n### Top 10 Countries of Origin ###\n ', top_country_of_origin)
      print('\n### Top 10 Filming Locations ###\n ', top_filming_location)

### Top 10 Genre ###
['Drama', 'Comedy', 'Action', 'Adventure', 'Thriller', 'Crime', 'Romance',
'Horror', 'Mystery', 'Sci-Fi']

### Top 15 Stars ###
['Dwayne Johnson', 'Mark Wahlberg', 'Ryan Reynolds', 'Matt Damon', 'Amy
Adams', 'Scarlett Johansson', 'Brad Pitt', 'Cate Blanchett', 'Samuel L Jackson',
'Liam Neeson', 'Tom Hanks', 'Jake Gyllenhaal', 'Colin Farrell', 'Gerard Butler',
'Christian Bale']

### Top 10 Directors ###
['Ridley Scott', 'Steven Spielberg', 'Antoine Fuqua', 'Clint Eastwood', 'Shawn
Levy', 'Michael Bay', 'Quentin Tarantino', 'Tim Burton', 'Zack Snyder', 'Robert
Rodriguez']

### Top 10 Countries of Origin ###
['United States', 'United Kingdom', 'Canada', 'France', 'Germany', 'China',
'Japan', 'Australia', 'Spain', 'Italy']

### Top 10 Filming Locations ###
['USA', 'Canada', 'UK', 'Australia', 'France', 'Spain', 'Italy', 'New
Zealand', 'Germany', 'Mexico']
```

```
[42]: df['Certificate'].value_counts()
```

```
[42]: R          945
      PG-13       727
      PG          229
      TV-MA        40
      G           22
      TV-14        12
      TV-PG         9
      NC-17         6
      TV-G          3
      TV-Y7         1
      Name: Certificate, dtype: int64
```

```
[43]: def measure_target(category, top_elements, df, target, measure):
      """
      Computes the target variable measured by the given category and top_
      ↪elements.
```


Parameters:

- category (str): The name of the category to measure by.*
- top_elements (list of str): The top elements to measure the target variable*
↳ for within the category.
- target (str): The name of the target variable to measure.*
- measure (str or function): The measure to use to aggregate the target*
↳ variable. This can be a string
 - specifying the name of the function to use (e.g., 'mean', 'median',*
↳ 'sum'), or a function object
 - that takes a Series and returns a scalar.*

Returns:

- dict of {str: float}: A dictionary where each key is an element and the*
↳ corresponding value is the
 - measured target variable for that element.*

```
'''
measure_target_top_elements = {}
for element in top_elements:
    measure_target = df.loc[df[category].str.contains(element), target].
↳agg(measure)
    # add every element's ROI, average, or median Profit rounding to the
↳nearest 100 thousands
    measure_target_top_elements[element] = measure_target #if not pd.
↳isna(measure_target) else -1
    # sort the dictionary in ascending order
    measure_target_top_elements = dict(sorted(measure_target_top_elements.
↳items(),
                                            key=lambda item: item[1],
↳reverse=True))
    return measure_target_top_elements
```

```
def plot_target_by_category(target, measure, categories=None,
↳top_elements_dict=None, y_metric=1e8, y_tick_upper=3.5):
    '''
    Plots a bar chart of a target variable measured by category, either for a
↳list of categorical variables or for the top elements
    within each category.

    Parameters:
    target (str): The name of the target variable to plot.
    measure (str): The type of measure to use for the target variable (e.g.
↳'median', 'mean', 'sum', 'count').
```

```

    categories (list of str): The names of the categories to plot. Either this
    ↪ or top_elements_dict must be provided.
    top_elements_dict (dict of {str: list of str}): A dictionary where each key
    ↪ is a category and the corresponding value
        is a list of the top elements to plot within that category. Either this
    ↪ or categories must be provided.
    y_metric (float): The scaling factor to use for the y-axis tick labels
    ↪ (default is 1e8).
    y_tick_upper (float): The upper limit of the y-axis tick labels (default is
    ↪ 3.5)
    '''

    # Validate input parameters
    if categories is None and top_elements_dict is None:
        raise ValueError('Either "categories" or "top_elements_by_category"
    ↪ argument must be provided.')
    if categories is not None and not isinstance(categories, list):
        raise ValueError('The "categories" argument must be a list.')
    if top_elements_dict is not None and not isinstance(top_elements_dict,
    ↪ dict):
        raise ValueError('The "top_elements_by_category" argument must be a
    ↪ dictionary.')
    if not callable(measure) and not isinstance(measure, str):
        raise ValueError('The "measure" argument must be ea string.')

    if top_elements_dict is not None:
        # Create a dictionary containing the category and a list of x and y
    ↪ values of the top elements
        data_to_plot = {}
        for idx, (category, top_elements) in enumerate(top_elements_dict.
    ↪ items()):
            # Compute the target variable measured by the given category and
    ↪ top elements
            target_measured_top_elements = measure_target(category,
    ↪ top_elements, df, target, measure)
            # Add the category and a tuple of the top elements and their
    ↪ corresponding target values to dict_xy
            x, y = list(target_measured_top_elements.keys()),
    ↪ list(target_measured_top_elements.values())
            data_to_plot[idx] = [category, (x, y)]

    else:
        # Create a dictionary containing the category and a list of x and y
    ↪ values
        data_to_plot = {}
        for idx, category in enumerate(categories):

```

```

        target_measured_by_category = df.groupby(category)[target].
↪agg(measure).sort_values(ascending=False)
        # get x and y axis values
        x, y = target_measured_by_category.index,
↪target_measured_by_category.values
        # add x,y as tuple to list
        data_to_plot[idx] = [category, (x, y)]

# get number of rows dynamically
num_plots = len(data_to_plot)
max_cols = 2

# Calculate number of rows and columns needed
num_rows = math.ceil(num_plots / max_cols)
num_cols = math.ceil(num_plots / num_rows)

# Create a 2x2 grid of subplots
fig, axs = plt.subplots(nrows=num_rows, ncols=num_cols, figsize=(8,
↪4*num_rows))

# Check if axs is an ndarray before flattening
if isinstance(axs, np.ndarray):
    axs = axs.flatten()
else:
    axs = [axs]

if num_plots % 2 == 1:
    axs[-1].axis('off')
    axs = axs[:-1] # Remove the last subplot

# axs[2,1].axis('off')

# Loop through each subplot and create a bar plot
for i, ax in enumerate(axs):

    # Create the bar plot
    sns.set_style('whitegrid')
    sns.barplot(x=data_to_plot[i][1][0], y=data_to_plot[i][1][1], data=df,
↪ax=ax, palette="Blues_r")

    # Set the x and y axis labels
    ax.set_xlabel(data_to_plot[i][0])
    ax.set_xticklabels(ax.get_xticklabels(), rotation=90, fontsize=9)
    ax.set_ylabel(f'{target} ({measure})')
    ax.set_yticks(ticks=np.arange(0, y_tick_upper, 0.5) * y_metric)
    ax.tick_params(axis='both', which='major', labels=9)

```

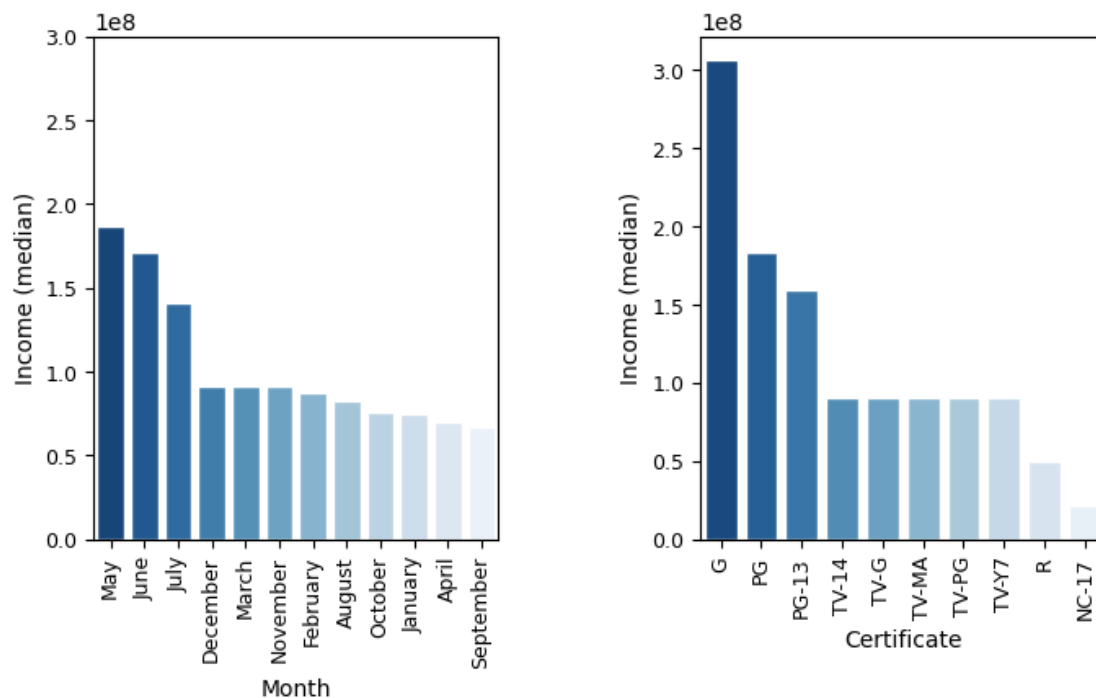
```
# Adjust the spacing between subplots
plt.subplots_adjust(wspace=0.5, hspace=0.75)

# Show the plot
plt.show()
```

The following plots are based on the median given the skewness of the distribution of the Income variable



```
[44]: categories=['Month', 'Certificate']

plot_target_by_category('Income', 'median', categories=categories,
↪ y_metric=1e8, y_tick_upper=3.5)
```





```
[45]: categories=['Month', 'Certificate']






plot_target_by_category('Profit', 'median', categories=categories,
↪ y_metric=1e8, y_tick_upper=2.5)
```





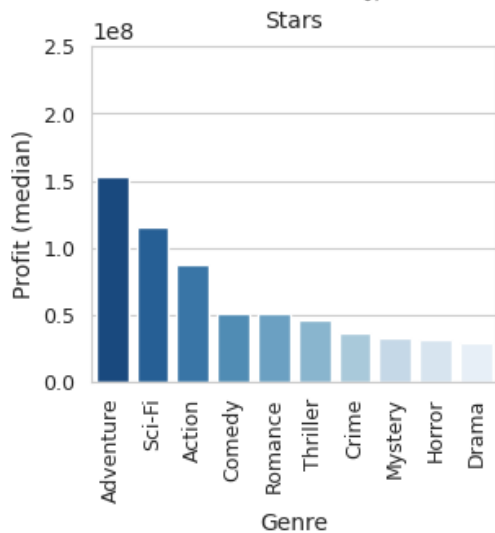


```
[46]: plot_target_by_category('ROI', 'median', categories=categories, y_metric=1,
    ↪ y_tick_upper=3)
```





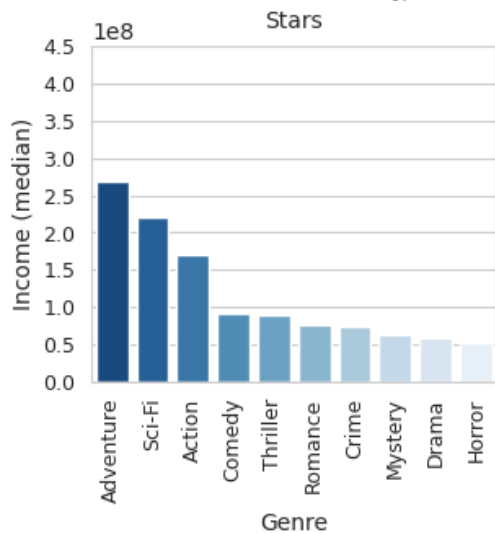


```
[47]: plot_target_by_category('ROI', 'median', top_elements_dict=top_elements_dict,
    ↪ y_metric=1, y_tick_upper=4)
```



```
[48]: plot_target_by_category('Profit', 'median',  
    ↪ top_elements_dict=top_elements_dict, y_metric=1e8, y_tick_upper=3)
```

```
[49]: plot_target_by_category('Income', 'median',  
    ↪ top_elements_dict=top_elements_dict, y_metric=1e8, y_tick_upper=5)
```



The exploratory data analysis provides multiple insights that can be utilized by movie production companies - the most important ones are: June, May, and July are the most profitable release months. The profitability is at least 50% higher in June and May than in all other months except July.

Adventure, Sci-Fi and Action are the ‘most profitable’ genres.

Dwayne Johnson, Samuel L. Jackson and Brad Pitt are the ‘most profitable’ actors to cast.

The rating of the movie is not significantly correlated with profit. Movies can be very profitable but not be relatively high rated by audiences.

Different production companies might want to pursue a different measure of financial performance. Maximizing ROI instead of profit could be more valuable to those. The following are a few of the data analysis insights: January is in the top 3 months if the aim is to maximize ROI, while for profit it’s only 8th highest.

Certificate NC-17 generates the highest ROI on average, while it’s least effective when the aim is profit.

Having Quentin Tarantino as a director suggests the highest chance at maximizing ROI, while he’s only 5th if the aim is profit.

Scarlett Johansson as a star generates the highest ROI on average, while only 11th highest when the aim is profit.

Adventure, Comedy and Romance are the top 3 genres for maximizing ROI

China, Spain and the United States are the top 3 filming locations for maximizing ROI.

```
[50]: def create_bins(column, percentiles, labels):

    # check whether the column is numeric or categorical
    if column in df.select_dtypes(exclude='number').columns:

        def extract_counts(column):
            # Split values in column by commas and create a series with a list
            ↪ of values for each row
            star_lists = df[column].str.split(', ')

            # Count the occurrences of each value in the series after exploding
            ↪ the lists
            star_counts = star_lists.explode().value_counts()

            # Calculate the total appearances of all values in each row's list
            ↪ and return as a series
            cast_counts = star_lists.apply(lambda stars: star_counts[stars].
            ↪ sum())
```

```

        return cast_counts

    # Extract the counts for the given column
    cast_counts = extract_counts(column)

    # compute the count cut offs based on the percentile cut offs
    cut_offs = cast_counts.quantile(percentiles)

    # make the function usable for numerical columns
    else:

        cut_offs = df[column].quantile(percentiles)

        cast_counts = df[column]

    # Create a list of bin edges based on the count cut offs
    bin_edges = [-float('inf')] + list(cut_offs) + [float('inf')]

    # Create bins using bin edges and labels
    bins = pd.cut(cast_counts,
                  bins=bin_edges,
                  labels=labels,
                  include_lowest=True,
                  duplicates='drop',
                  ordered=False)

    # Add binned column to the original dataframe
    df[f'{column}_class'] = bins

```

```

[51]: # create the new columns and define the categories
create_bins('Stars', [0.25, 0.5, 0.75, 0.95], ['unknown', 'known',
↪ 'well_known', 'famous', 'iconic'])
create_bins('Directors', [0.5, 0.75, 0.97], ['unknown', 'known', 'well_known',
↪ 'famous'])
create_bins('Budget', [0.5, 0.75], ['low', 'moderate', 'high'])
create_bins('Runtime', [0.5, 0.95], ['short', 'moderate', 'long'])

```

5.0.1 Inspect new column Stars_class & Directors_class

```

[52]: # inspect the new columns
sns.boxplot(data=df, x='Stars_class', y='Profit')
df[['Profit', 'Stars_class']].groupby('Stars_class').describe()

```

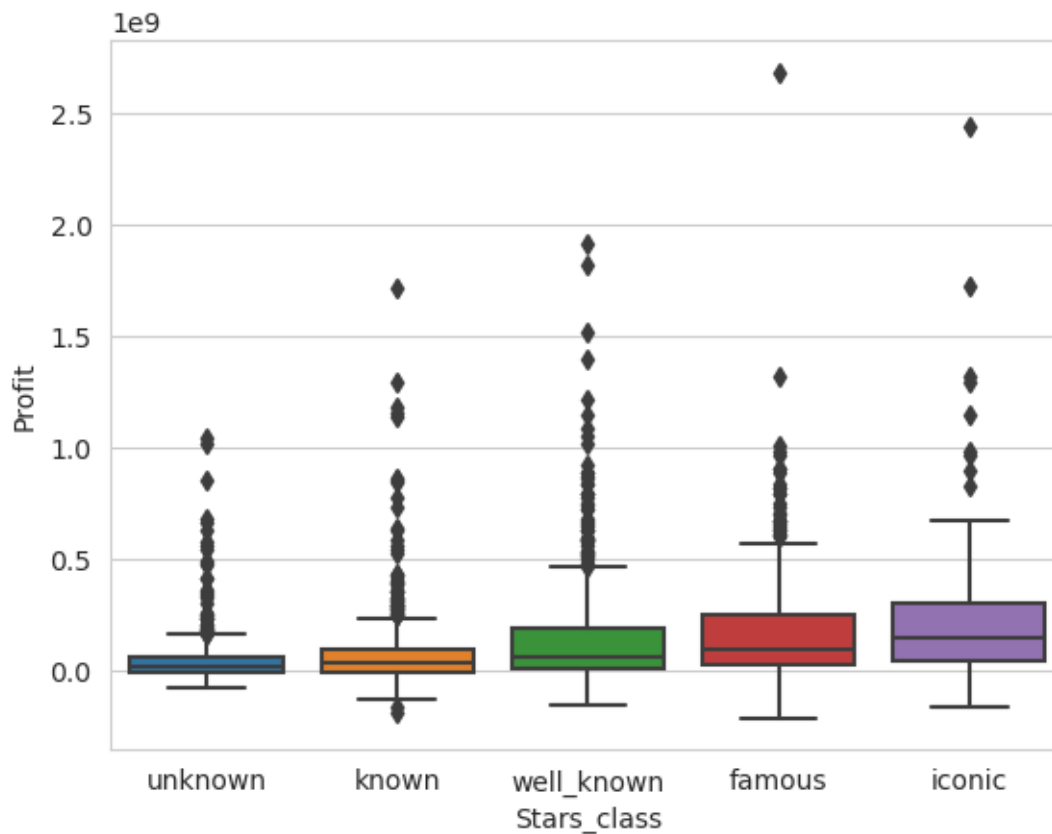
```

[52]:
      Profit
count      mean      std      min      25% \
Stars_class

```

unknown	510.0	5.209216e+07	1.231403e+08	-70000000.0	-5000000.0
known	527.0	8.548577e+07	1.788172e+08	-186000000.0	-2000000.0
well_known	481.0	1.646549e+08	2.612742e+08	-150000000.0	13000000.0
famous	388.0	1.839330e+08	2.579545e+08	-209000000.0	29000000.0
iconic	88.0	2.825227e+08	4.129427e+08	-160000000.0	43750000.0

	50%	75%	max
Stars_class			
unknown	22000000.0	62750000.0	1.043000e+09
known	42000000.0	99000000.0	1.717000e+09
well_known	67000000.0	196000000.0	1.918000e+09
famous	98000000.0	249250000.0	2.686000e+09
iconic	151000000.0	301750000.0	2.442000e+09

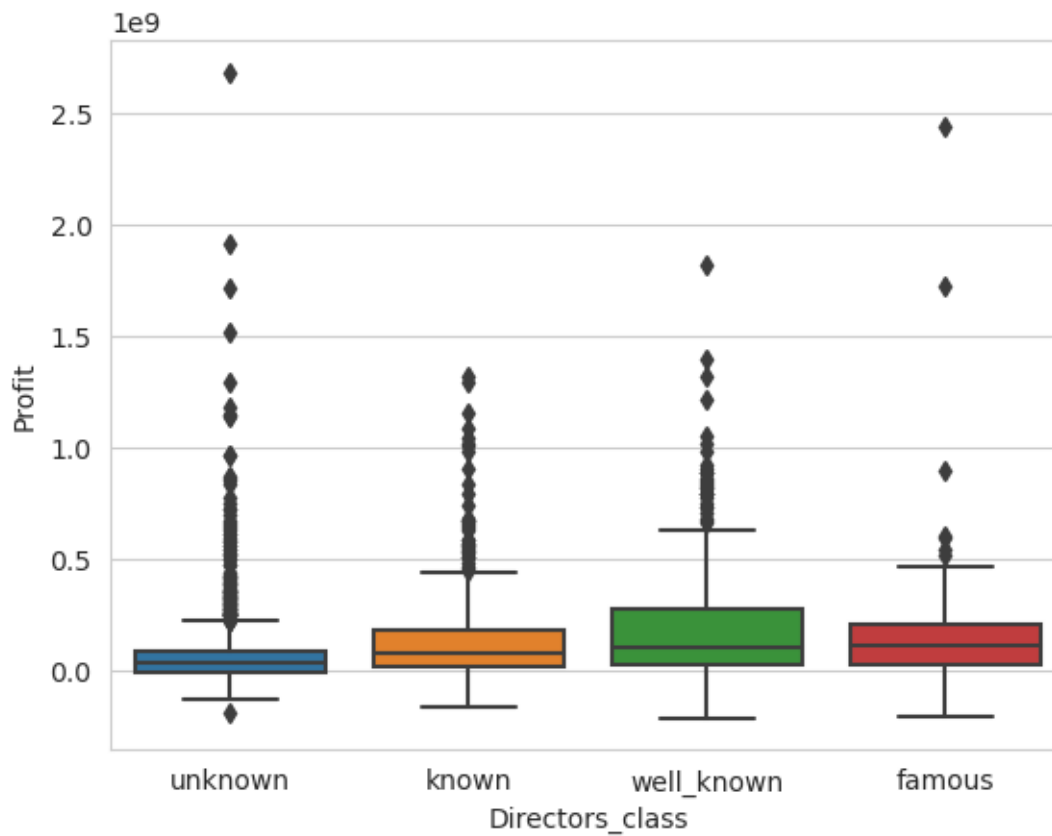


```
[53]: # inspect the new columns
sns.boxplot(data=df, x='Directors_class', y='Profit')
df[['Profit', 'Directors_class']].groupby('Directors_class').describe()
```

[53]:

	Profit				
	count	mean	std	min	25%
Directors_class					
unknown	1204.0	8.305897e+07	1.904146e+08	-186000000.0	-2000000.0
known	348.0	1.602471e+08	2.325630e+08	-160000000.0	16250000.0
well_known	397.0	2.023904e+08	2.681181e+08	-209000000.0	32000000.0
famous	45.0	2.428222e+08	4.598986e+08	-200000000.0	29000000.0

	50%	75%	max
Directors_class			
unknown	35000000.0	89250000.0	2.686000e+09
known	81000000.0	187250000.0	1.319000e+09
well_known	103000000.0	277000000.0	1.825000e+09
famous	119000000.0	214000000.0	2.442000e+09




```
[54]: sns.boxplot(data=df, x='Budget_class', y='Profit')
df[['Profit', 'Budget_class']].groupby('Budget_class').describe()
```

[54]:

		Profit				
	count	mean	std	min	25%	
Budget_class						
low	1199.0	4.429358e+07	7.806267e+07	-40000000.0	-2000000.0	
moderate	309.0	1.103625e+08	1.519539e+08	-74000000.0	20000000.0	
high	486.0	3.288786e+08	3.610488e+08	-209000000.0	85750000.0	

		50%	75%	max	
Budget_class					
low	28000000.0	66000000.0	6.670000e+08		
moderate	73000000.0	147000000.0	1.085000e+09		
high	227500000.0	499000000.0	2.686000e+09		



[55]:


```
sns.boxplot(data=df, x='Runtime_class', y='Profit')
df[['Profit', 'Runtime_class']].groupby('Runtime_class').describe()
```

[55]:





		Profit				
	count	mean	std	min	25%	
Runtime_class						

short	1009.0	8.430327e+07	1.520224e+08	-150000000.0	0.0
moderate	888.0	1.515822e+08	2.514673e+08	-200000000.0	10000000.0
long	97.0	2.822474e+08	4.746611e+08	-209000000.0	7000000.0





	50%	75%	max
Runtime_class			
short	47000000.0	103000000.0	1.300000e+09
moderate	58000000.0	185250000.0	1.825000e+09
long	94000000.0	381000000.0	2.686000e+09



```
[56]: plot_target_by_category('Profit', 'median', categories=['Directors_class',
↳ 'Stars_class', 'Budget_class', 'Runtime_class'], y_metric=1e8,
↳ y_tick_upper=2.5)
```



```
[57]: plot_target_by_category('ROI', 'median', categories=['Directors_class', 'Stars_class', 'Budget_class', 'Runtime_class'], y_metric=1, y_tick_upper=3)
```



Further Business Insights

Movies tend to be more profitable with higher-profile actors in them and less profitable with lower-profile actors in them.

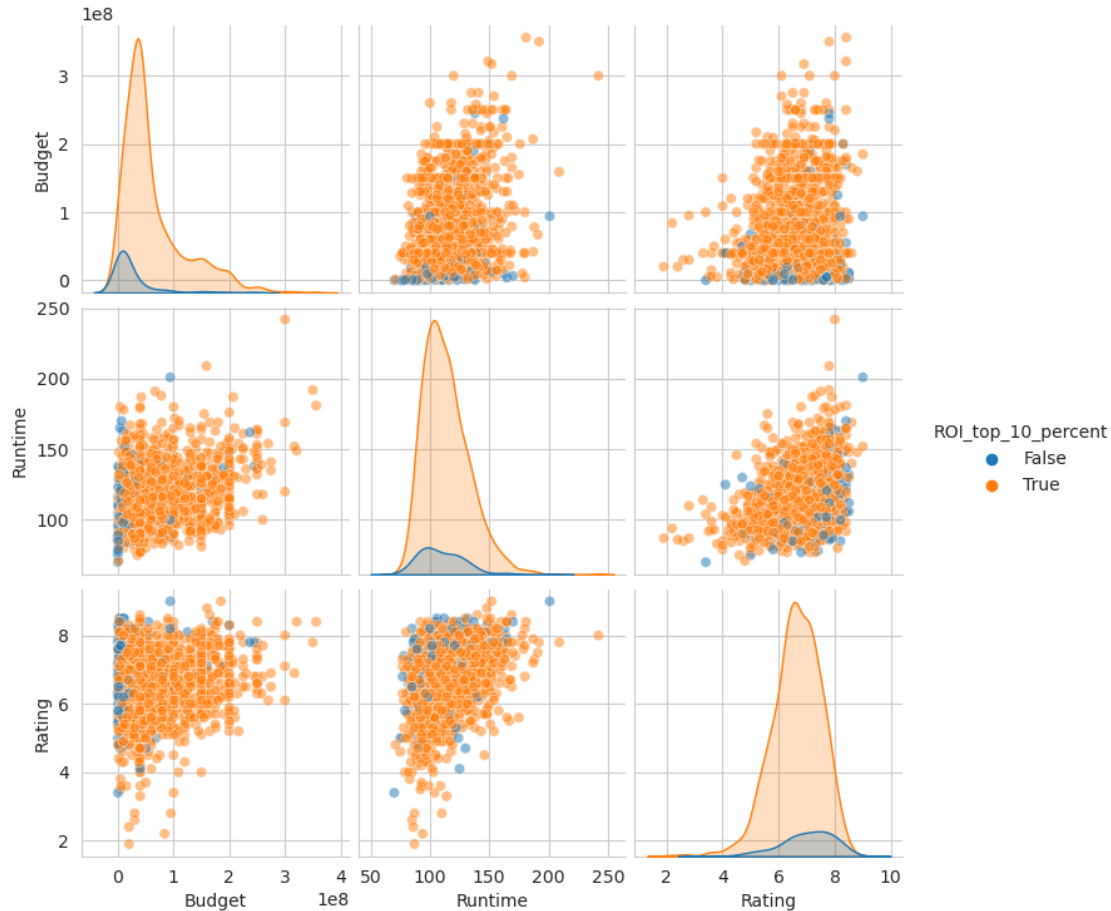
Longer movies tend to be more profitable than shorter movies.

Movies without age restriction tend to be more profitable than movies with other certificates.

```
[58]: # Create a new column to identify if ROI is below threshold
df['ROI_bottom_25_percent'] = df['ROI'] < df['ROI'].quantile(0.25)
df['ROI_top_10_percent'] = df['ROI'] < df['ROI'].quantile(0.90)
```

```
# Select variables to plot
numerics_to_plot = ['Budget', 'Runtime', 'Rating']
categoricals_to_plot = ['Stars_class', 'Directors_class', 'Month']

# Create scatter plot matrix using seaborn with df_below_threshold as hue
sns.pairplot(df, vars=numerics_to_plot, hue='ROI_top_10_percent',
            plot_kws={'alpha':0.5})
plt.show()
```



Further interesting finding:

Although Rating is not that important for our analysis because that variable is unknown before the production, it is interesting to notice that it seems as though a low budget regardless of Rating tends to yield a lower Return on Investment

5.1 Practical Implications

Release your movie in May, June or July

The impact of user ratings is not significant for profitability. So, don't be scared from trying

something new even if you risk lower ratings.

Adventure, sci-fi or action movies are profitable genre

Having The Rock, Samuel L Jackson, or Brad Pitt in your movie would be a homerun

For smaller production companies who focus more on ROI than on profit, the optimal choices can differ and should be chosen accordingly

6 Income Prediction

Predicting income can be insightful when evaluating a movie's financial success because it gives an estimate of the total revenue generated by the movie. While profit and ROI are important metrics for measuring financial success, they are dependent on multiple factors beyond just the revenue generated by the movie. For example, profit depends on the movie's budget, marketing costs, distribution fees, and other expenses.

On the other hand, predicting income provides a straightforward estimate of the revenue generated by the movie, which can be compared to the movie's budget to determine whether it was financially successful. Additionally, predicting income can help in forecasting the potential revenue that a movie might generate in the future.

Overall, predicting income from a movie database can be a useful tool in evaluating a movie's financial success as it provides a clear picture of the movie's revenue generating potential.

6.1 Feature Engineering

Before the models will be trained 'Dummy Variables' are created for 'Categorical Variables'. Furthermore, highly correlated features are excluded (e.g. certain filming locations) and the numerical variables with skewed distributions, including the target variable were log transformed. The dataset is split into 80% training and 20% test data to later evaluate the performance of the model.

```
[59]: df_model = df.copy()
```

```
[60]: # features that have multiple entries in each cell
columns = ['Genre', 'Country_of_origin', 'Certificate']

# function for most of the remaining columns that need to be changed
for col in columns:
    tmp_df = df_model[col].str.get_dummies(sep=', ').add_prefix(f'{col}_')
    # Drop the first column to interpret coefficients
    tmp_df = tmp_df.iloc[:, 1:] # Drop the first column
    df_model = pd.concat([df_model, tmp_df], axis=1)
    df_model = df_model.drop(col, axis=1)
```

```
[61]: for col in ['Stars_class', 'Directors_class', 'Month', 'Filming_location']:
    df_model = pd.get_dummies(df_model, columns=[col], drop_first=True)
```

```
[62]: to_drop = [ 'Directors', 'Stars', 'Profit', 'ROI',
                'Rating', 'ROI_top_10_percent', 'ROI_top_10_percent',
                'ROI_bottom_25_percent', 'Budget_class', 'Runtime_class']
df_model = df_model.drop(to_drop, axis=1)
```

```
[63]: from sklearn.preprocessing import FunctionTransformer

# Separate the target variable and input features
X = df_model.drop('Income', axis=1)
y = df_model['Income']

# Define a log transformation function
log_transform = FunctionTransformer(np.log1p)

# Apply the log transformation to the X data
X['Budget'] = log_transform.transform(X['Budget'])
X['Runtime'] = log_transform.transform(X['Runtime'])

# Apply the log transformation to the y data
y = log_transform.transform(y)
```

```
[64]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
```

```
[65]: # Get correlation between features
corr = X_train.corr()
#sns.heatmap(corr, cmap='coolwarm', annot=True)
#plt.show()

# Potentially drop highly correlated features
cor_matrix = X_train.corr().abs()
upper_tri = cor_matrix.where(np.triu(np.ones(cor_matrix.shape), k=1).astype(np.
    ↪bool_))
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.
    ↪75)]
X_train = X_train.drop(to_drop, axis=1)
X_test = X_test.drop(to_drop, axis=1)

print('Highly correlated feature(s) dropped: ', to_drop)
```

Highly correlated feature(s) dropped: ['Country_of_origin_Lebanon',
 'Country_of_origin_Qatar', 'Country_of_origin_Uruguay', 'Filming_location_Iran',
 'Filming_location_Kenya', 'Filming_location_Lebanon',

```
'Filming_location_Slovenia', 'Filming_location_South Korea']
```

6.1.1 Required Imports, Functions, and Tuning parameters

```
[66]: from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error,
    ↪mean_absolute_percentage_error, median_absolute_error
from sklearn import tree

def cv_optimize(clf, parameters, X, y, n_jobs=2, n_folds=5, score_func=None):
    if score_func:
        gs = GridSearchCV(clf, param_grid=parameters, cv=n_folds,
    ↪n_jobs=n_jobs, scoring=score_func)
    else:
        gs = GridSearchCV(clf, param_grid=parameters, n_jobs=n_jobs, cv=n_folds)
    gs.fit(X, y)
    print('BEST Parameters: ', gs.best_params_)
    # print(f'BEST cross-validated {score_func}: ', gs.best_score_)
    # print(gs.grid_scores_)
    best = gs.best_estimator_
    return best

def median_absolute_percentage_error(y_test, y_pred):
    return np.median(np.abs((y_test - y_pred) / y_test)) * 100

def print_metrics(clf, X_test, y_test, Y_pred, y_train):

    # transform back to original scale to interpret mean error
    Y_pred_exp = np.expm1(Y_pred)
    y_test_exp = np.expm1(y_test)
    y_train_exp = np.expm1(y_train)

    # R2 scores on the log transformed data to stabilize the variance due to
    ↪skewed data
    # for linear regression
    print('Train R2-score: ', format(clf.score(X_train, y_train.values.
    ↪ravel()), '.2'))
    print('Test R2-score: ', format(clf.score(X_test, y_test.values.ravel()), '.
    ↪2'))
    r2 = r2_score(y_test.values.ravel(), Y_pred)

    # medape = median_absolute_percentage_error(y_test, Y_pred)
```

```

#medae = median_absolute_error(y_test.values.ravel(), Y_pred)
mae = mean_absolute_error(y_test_exp.values.ravel(), Y_pred_exp)
# mse = mean_squared_error(y_test.values.ravel(), Y_pred)
mape = mean_absolute_percentage_error(y_test_exp.values.ravel(), Y_pred_exp)

# Print metrics
print('Mean Absolute Error:', round(mae, 2))
#print('Median Absolute Error:', round(medae, 2))
print('R-squared:', round(r2, 2))
#print('MEDAPE: {:.2f}%'.format(medape))
print('MAPE: {:.2f}%'.format(mape))

def plot_tree(clf, features):

    # Setting dpi = 300 to make image clearer than default
    fig, axes = plt.subplots(nrows=1, ncols=1, figsize=(60,60), dpi=300)

    tree.plot_tree(clf,
                    feature_names=features,
                    # string class names must equate to their numerical representation
    ↪ in ascending order
                    class_names=['No', 'Yes'],
                    filled=True,
                    fontsize=40)

lr_params = {
    'fit_intercept': [True, False],
    'copy_X': [True, False],
    'n_jobs': [-1, 1, 2]
}

dt_params = {
    'max_depth': [2, 4, 6, 8],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [2, 4, 7],
    'random_state': [0, 42, 100, 200]
}

rf_params = {
    'n_estimators': [100, 200],
    'max_depth': [2, 4, 5, 8],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [4, 12, 16],
    'random_state': [0, 42, 100, 200]
}

```


6.2 OLS Regression

Investigating variables with significant p-values serves to analyze their importance for predicting income. Their coefficients indicate whether those variables have negative or positive impact on the target variable. Due to the log transformation the scale of the variables may not be absolute, but the coefficients can still be used to understand the direction and magnitude of their impact on the target variable.

```
[67]: import statsmodels.api as sm

# Add a constant column to X_train
X_train_c = sm.add_constant(X_train)

# Fit the OLS model using statsmodels
model = sm.OLS(y_train, X_train_c)
results = model.fit()

# Get the p-values for the coefficients
p_values = results.pvalues[1:] # Exclude the constant term

# Print the significant variables along with their p-values
significant_vars = p_values[p_values < 0.05]

coef = results.params[1:]

# Print significant variables (p-value < 0.05)
print(coef[p_values < 0.05])
```

Year	-0.045142
Runtime	2.439385
Budget	0.252611
Genre_Animation	0.679508
Genre_Drama	-0.682251
Genre_Romance	0.299828
Genre_Western	-1.292025
Country_of_origin_Australia	-0.659255
Country_of_origin_Belgium	-1.044146
Country_of_origin_China	0.624504
Country_of_origin_Cyprus	4.406649
Country_of_origin_Czech Republic	-1.084357
Country_of_origin_Israel	-4.429202
Country_of_origin_Serbia	-8.349648
Country_of_origin_United States	1.213661
Certificate_R	-1.236016
Stars_class_well_known	0.513749
Stars_class_famous	0.675317
Stars_class_iconic	0.730054
Directors_class_known	0.287000

```
Directors_class_well_known      0.475429
Filming_location_Argentina      2.929810
dtype: float64
```

7 Models

The model selection is based on the business objective, which is to provide insights into the factors that drive a movies financial performance. Interpretable models can provide clear insights into these factors. Moreover, performance metrics, such as mean absolute error, MAPE, and R squared server to evaluate the model.

7.1 Model Linear Regression

```
[68]: # Create and tune a linear regression model
lr = cv_optimize(LinearRegression(), lr_params, X_train, y_train,
    ↪score_func="r2")

# fit the model
lr.fit(X_train, y_train)

# Use the model to make predictions on the test set
Y_pred = lr.predict(X_test)

print_metrics(lr, X_test, y_test, Y_pred, y_train)
```

```
BEST Paramaters: {'copy_X': True, 'fit_intercept': False, 'n_jobs': -1}
Train R2-score:  0.47
Test R2-score:   0.32
Mean Absolute Error: 119787696.01
R-squared: 0.32
MAPE: 41.97%
```

7.2 Model Decision Tree

```
[69]: # Create and tune a DecisionTreeRegressor
dt = cv_optimize(DecisionTreeRegressor(), dt_params, X_train, y_train,
    ↪score_func="r2")

# fit the model
dt.fit(X_train, y_train)

# Use the model to make predictions on the test set
Y_pred = dt.predict(X_test)

# show results
print_metrics(dt, X_test, y_test, Y_pred, y_train)
```

```
# show the decision tree
# plot_tree(dt, X_train.columns.tolist())
```

BEST Paramaters: {'max_depth': 4, 'min_samples_leaf': 4, 'min_samples_split': 10, 'random_state': 42}
 Train R2-score: 0.42
 Test R2-score: 0.37
 Mean Absolute Error: 101574322.81
 R-squared: 0.37
 MAPE: 38.67%

7.3 Model Random Forest

```
[70]: # Create and tune a RandomForestRegressor
rf = cv_optimize(RandomForestRegressor(), rf_params, X_train, y_train,
    ↪ score_func="r2")

# fit the model
rf.fit(X_train, y_train)

# Use the model to make predictions on the test set
Y_pred = rf.predict(X_test)

# show results
print_metrics(rf, X_test, y_test, Y_pred, y_train)
```

BEST Paramaters: {'max_depth': 8, 'min_samples_leaf': 16, 'min_samples_split': 2, 'n_estimators': 200, 'random_state': 100}
 Train R2-score: 0.53
 Test R2-score: 0.47
 Mean Absolute Error: 88142418.19
 R-squared: 0.47
 MAPE: 30.06%


Feature Importances by Random Forest

```
[71]: # plot top features importances determined by RandomForest
features = X_train.columns
importances = rf.feature_importances_
indices = np.argsort(importances)[-10:] # get the top 10 feature indices

f = plt.figure()
f.set_figwidth(5)
f.set_figheight(5)
plt.title('Top 10 Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='darkblue',
    ↪ align='center')
plt.yticks(range(len(indices)), features[indices])
```

```
plt.xlabel('Relative Importance')
```

```
[71]: Text(0.5, 0, 'Relative Importance')
```



8 Analysis of the Models

Analyzing Coefficients, features importances as indicated by the Random Forest and looking at the decision Tree, the following features are important for a movies' financial performance: Runtime, Budget, Certificate, Certificate_TV-PG, Stars_class, Filming_location_Serbi, Country_of_origin_USA, and different Genre, including Drama, Western, Animation, and Adventure. These insights also inspired going into some more depth in the Data Analysis insights that are mentioned above.

8.0.1 Practical Implications

As moderate budget movies do not seem to achieve much higher profits than low budget productions, it is not worth to have a medium budget movie, either try to keep the budget low or go big.

The profitability of a movie seems to be increasing with higher profile actors in them, so logically casting more famous actors promises more financial success.

Make movies longer and try to not receive an age restriction

The best performing model explains about 50% of the variance in the dependent variable, which means that there is still a huge potential to uncover more factors that contribute towards the financial performance of a movie.

8.0.2 Limitations

As the best r^2 score suggests, there are many other factors influencing the financial performance of these movies. Therefore, the insights should be taken as absolutes. Moreover, the analysis is based on correlation not causation, a fact often being overshadowed.

The decision Tree is not transformed back to the original scale. So as of now it is only useful for extracting the features that are split on, but not what the cut-offs are for Budget and Runtime.

I might be interesting to further explore the movie titles, which have been left out for this analysis.

[]: